



www.veritglobal.com

# Tech. Blueprint

## Contents:

1. Interface patterns & recommendations (with ranking, decision tree, anti-patterns)
2. Architecture: **12 core** modules + **4 context** boxes, ASCII diagram
3. Data model & repositories (per module + cross-cutting contracts)
4. How to run (K8s, scale units, DR/multi-region, SLOs/alerts, security)
5. Implementation plan & checklists (MVP → GA)
6. Failure modes & runbooks
7. Glossary & appendices (reason codes, payload schemas, API sketch)

## 1) Interface patterns & recommendations

### 1.1 The core invariant (applies to every interface)

- Compute payouts **deterministically**, seal a **transcript** and **output digest**, then **release funds only** if **replay digest == transcript digest** *and* an **acceptance bundle** (ACK/CT/optional SPV) meets **freshness** and **quorum**.

### 1.2 Options (8)

#	Interface pattern	Best for	Pros	Cons
1	<b>One-call Authorization Gate (REST/gRPC)</b> + tiny verifiers	Quick overlay, minimal change	Simple; reason-coded holds; cryptographic audit	Add 2–3 verifier webhooks; pass IDs through rails
2	<b>Batch flat-file (Tier-0)</b> + attestations	Weekly/biweekly; low lift	Fastest start; no streaming	Windowed latency; later upgrade to streaming

#	Interface pattern	Best for	Pros	Cons
3	<b>ERP-native adaptor</b> (e.g., NetSuite)	Finance-led, strict audit	Tight audit trail; clean GL mapping	Some scripting/config; chunk at volume
4	<b>SDK-first (Stripe-style)</b>	Developer-led builds	Great DX; feature flags	Manage keys/versions; still need attestations
5	<b>Event streaming</b> (Kafka/Pub/Sub)	High volume, near real-time	Smooth scale; predictable close	More platform work than files
6	<b>PSP plugin + SPV receipts</b>	Keep existing rails	Airtight reconciliation	Wire receipt pickup within freshness
7	<b>On-/cross-chain SPV adaptor</b>	On-ledger touchpoints	Deterministic acceptance across domains	Finality/headers to monitor
8	<b>Ops copilot</b> (service/agent)	Large ops teams	Fewer escalations; auto-nudges	Support layer, not primary integration

**Default ranking (typical buyer):** 1 → 6 → 2 → 5 → 3 → 4 → 8 → 7

### 1.3 Decision tree (text)

```

Start → Need value in weeks (Y) → Choose 1 + 2; add 6 later
      |
      N → Tight SLA or heavy scale (Y) → 1 + 5 (+6)
      |
      N → Finance-led & ERP-centric (Y) → 3 + 1 (+6)
      |
      N → Product/dev-led (Y) → 4 + 1 (+6)
      |
      Crypto/on-ledger (Y) → add 7 to acceptance matrix

```

### 1.4 Anti-patterns to avoid

- Becoming the **payer of record** (replace rails) → prefer overlay + SPV.
  - **Per-payment** synchronous gating → gate **per window**.
  - Re-implementing compute in ERP → keep ERP as **system of record**, not allocation engine.
  - Skipping evidence (no ACK/CT/SPV) → defeats the purpose.
  - Dropping determinism guardrails (unordered data, FP math, wall-clock) → breaks replay.
-

## 2.1 Big picture ASCII (copy-safe)

*Process:* compile policy → IR; fix window cadence/timezone; shard function/version; rounding scale S; acceptance F/Q; checkpoint cadence; sign/version.

*Repositories:* Git/config; Manifest table (OLTP); artifact bucket.

*Flows:* **Gets** Governance [C9], Security [C2] → **Sends** params to [C4][C5][C7].

## **[C2] Security & Keys**

*Interfaces:* KMS/HSM, cert issuance, signing API.

*Process:* mTLS everywhere; sign transcript roots/manifests; rotate keys; audit.

*Repositories:* KMS/HSM; key metadata (OLTP).

*Flows:* **Gets** Governance [C9] → **Sends** signatures to [C6] and [C1].

## **[C3] Ingestion & Normalizer**

*Interfaces:* Event API; file intake; streams.

*Process:* schema validate; derive/validate `window_id`; idempotency; assign `partition_id`; canonical timestamps.

*Repositories:* raw/normalized lake; topics; idempotency index.

*Flows:* **Gets** External [CTX-1], mappings [C12] → **Sends** canonical events to [C4].

## **[C4] Partitioned Logs & Window Manager**

*Interfaces:* append per (tenant, window, partition); watermark queries.

*Process:* single-writer; lexicographic fold; monotone watermark to close.

*Repositories:* compacted logs (open); window state (OLTP); archived logs.

*Flows:* **Gets** [C3], rules [C1] → **Sends** close + ordered stream to [C5].

## **[C5] Deterministic Execution Engine**

*Interfaces:* window close trigger; replay versioning.

*Process:* i128 accumulation; late quantization; deterministic carry; no FP/wall-clock/hidden I/O.

*Repositories:* worker KV; checkpoints; execution index.

*Flows:* **Gets** [C4], [C1] → **Sends** allocation records to [C6].

## **[C6] Transcripts & Digests (tiered)**

*Interfaces:* get transcript/output digest/segment list.

*Process:* content-addressed segments (inputs/checkpoints/outputs); compute `output_digest` + trailer (fold order, watermark, manifest hash); sign/publish `transcript_root`.

*Repositories:* object store tree; transcript index (OLTP).

*Flows:* **Gets** from [C5], signatures [C2] → **Sends** digest/root to [C7][C8][CTX-2][C11].

## **[C7] Verifiers (+ Acceptance Matrix)**

*Interfaces:* proof intake: **ACK** (finance), **CT** (tax/KYC/rights), **SPV** (provider receipts or

on-chain headers); acceptance query.

*Process:* validate; enforce freshness F & quorum Q; reason-coded outcomes; bind acceptance to transcript/output\_digest.

*Repositories:* acceptance (OLTP); proofs (object store); cache for not-expired.

*Flows:* **Gets** digest/root [C6], external proofs [CTX-1/3] → **Sends** acceptance to [CTX-2]; snapshot to [C10].

## **[C8] Replay & Equivalence**

*Interfaces:* start replay; get equality/mismatch.

*Process:* reproduce from transcripts only; block on mismatch; store diffs.

*Repositories:* job log (OLTP); replay logs (object store).

*Flows:* **Gets** [C6] → **Sends** result to [CTX-2], [C9], [C10].

## **[C9] Governance & Safe-Change**

*Interfaces:* propose/promote/rollback; canary cohorts; re-shard dual-write.

*Process:* N-window success before promotion; rollback on violations; INVALID + checkpoint recovery.

*Repositories:* governance ledger (append-only, signed); manifest snapshots.

*Flows:* **Gets** [C7][C8] → **Sends** activations to [C1]; rules to [C4][C6]; notifications to [C10].

## **[C10] Observability & KPIs**

*Interfaces:* metrics/logs/traces; dashboards/alerts.

*Process:* time-to-release; replay-equality; reason-codes; carry remainder; watermark lag; proof freshness.

*Repositories:* metrics TSDB; logs; traces; BI warehouse.

*Flows:* **Gets** [C4–C8], [CTX-3] → **Sends** alerts; feeds [C9].

## **[C11] Reconciliation & Provider Adapters**

*Interfaces:* provider report intake; normalize; post SPV; ERP export mappers.

*Process:* map `provider_batch_id` ↔ {`window_id`, `output_digest`}; reconcile totals; raise diffs as reasons; feed SPV back to Verifiers.

*Repositories:* raw drop-zone; normalized receipts (OLTP); recon dashboards.

*Flows:* **Gets** from [CTX-3] → **Sends** SPV to [C7]; GL maps to [C12]; metrics to [C10].

## **[C12] Connector Layer to ERP/GL/PSP**

*Interfaces:* field mappings; webhook inbox; export jobs; optional ERP `eligible_to_pay` flag.

*Process:* stamp `window_id`, `output_digest`, `transcript_url`, `provider_batch_id` on bills/payments; filter ALLOW for pay-runs.

*Repositories:* connector state (OLTP); payload archives.

*Flows:* **Gets** decisions [CTX-2], recon [C11] → **Sends** to External [CTX-1]; optional echoes to [C3].

### [CTX-1] External Systems

*Interfaces:* file drops, streams, read-only DB views.

*Process:* land/validate/promote; supply KYC/Tax proofs; return provider receipts.

*Repositories:* staging buckets/queues.

*Flows:* **Sends** events to [C3]; proofs to [C7]; receipts to [C11].

### [CTX-2] Authorization Gate (window-level façade)

*Interfaces:* authorize `window_id` → ALLOW/HOLD with reasons per principal.

*Process:* require (replay digest == transcript digest)  $\wedge$  (acceptance F/Q met).

*Repositories:* authorization log (OLTP).

*Flows:* **Gets** digest [C6], acceptance [C7], replay [C8] → **Sends** ALLOW to [CTX-3]; eligible refs to [C12].

### [CTX-3] Disbursement via PSP/Bank

*Interfaces:* originate batches; tag (`window_id`, `output_digest`); return `provider_batch_id` and reports.

*Process:* pay ALLOW; keep batch/report artifacts for SPV.

*Repositories:* payout batch registry (OLTP); raw provider reports.

*Flows:* **Gets** ALLOW [CTX-2] → **Sends** metadata to [C11]; receipts to [C7] via [C11].

### [CTX-4] Acceptance Matrix (Ops View)

*Interfaces:* read-only per-window matrix (kinds, freshness, quorum, expiry, status).

*Process:* materialized view from [C7]; no business logic.

*Repositories:* materialized view + cache.

*Flows:* **Gets** from [C7]; **Sends** none.

---

## 3) Data model & repositories

### 3.1 Cross-cutting IDs

- `window_id`, `tenant_id`, `partition_id`, `policy_version`, `output_digest`,  
`provider_batch_id`, `transcript_root`.

### 3.2 Canonical contracts (sketch)

- **EventRecord (C3):** {tenant\_id, window\_id, event\_id, ts\_logical, bucket\_id, principal\_id, currency, amount\_minor, payload\_digest, provider\_id}
- **AllocationRecord + Trailer (C6):** allocations[] + {watermark, fold\_order\_desc, policy\_manifest\_hash}
- **PayoutHeader (C6/C7):** {window\_id, policy\_version, output\_digest, acceptance\_requirements{F,Q,kinds}}
- **AcceptanceRecord (C7):** {window\_id, kind ∈ {ACK,CT,SPV}, subject\_id, status, quorum, expires\_at, reason\_code, signature}
- **AuthDecision (CTX-2):** {window\_id, principal\_id, output\_digest, decision ∈ {ALLOW,HOLD}, reason\_code}

### 3.3 Repositories by module (authoritative store)

- **C1** Git + Manifest (OLTP); artifacts in object store
- **C2** KMS/HSM; key metadata (OLTP)
- **C3** Object store (raw/normalized), Topics, Idempotency (OLTP)
- **C4** Topics (open), Window state (OLTP), Archived logs
- **C5** Worker KV + Checkpoints (object store)
- **C6** Object store tree + Transcript index (OLTP)
- **C7** Acceptance (OLTP) + Proofs (object store)
- **C8** Replay jobs (OLTP) + logs (object store)
- **C9** Governance ledger (OLTP, append-only) + snapshots
- **C10** Metrics/Logs/Traces backends + Warehouse
- **C11** Raw receipts (object store) + Normalized receipts (OLTP)
- **C12** Connector state (OLTP) + archives
- **CTX-2** Auth log (OLTP); **CTX-3** Payout registry (OLTP); **CTX-1** staging buckets; **CTX-4** materialized view

### 3.4 Security & retention

- **Transcripts** immutable; 7–10 years; WORM/immutability.
- **Proofs** expire at `expires_at`, but **decisions**/audit retained long-term.
- Keep **PII out of transcripts**; store sensitive CT/KYC evidence separately with KMS keys.

---

## 4) How to run (platform & scale)

### 4.1 Deployment model

- **Default:** Multi-AZ SaaS single region → Active-passive DR → Active-active (tenant-pinned windows).
- **Alt:** Private cloud/on-prem: same shapes, self-hosted services.

## 4.2 Kubernetes layout

- **Namespaces per env** (dev/stage/prod). HPA on [C3][C4][C5][C7].
- API Gateway/Ingress + mTLS + NetworkPolicies.
- Secrets via KMS; short-lived tokens (SPIFFE/SPIRE optional).

## 4.3 State & queues

- **Kafka/PubSub** for [C3][C4]; 3→9 brokers, ISR across AZs.
- **Postgres HA** for [C1][C7][C9][C12][CTX-2]; shard later (Citius/Hyperscale).
- **Object store** for [C6][C11]; cross-region replication + Object Lock.
- **Redis** optional for caches/webhook dedupe.
- **Observability:** Prometheus + ELK/OpenSearch + Jaeger/Tempo; Warehouse for BI.

## 4.4 Scale units & autoscaling

- Partition count  $\kappa$  by window (target 2–5M events/partition); one **engine worker per (partition, window)**.
- Autoscale triggers: **queue depth** (C3), **watermark lag/open windows** (C4), **active partitions & CPU** (C5), **proof backlog & expiry proximity** (C7).
- Backpressure: 429 beyond tenant quota; pause promotion if `time-to-release` SLO degrades.

## 4.5 Multi-region

- **Phase 1:** single region multi-AZ.
- **Phase 2:** active-passive DR (RPO  $\leq$  5m, RTO  $\leq$  30m).
- **Phase 3:** active-active via tenant pinning or independent windows; reconcile at [C11].

## 4.6 Boot sequence

1. Provision platform (Kafka, Postgres(HA), Object store, Redis, Observability).
2. Bring up control plane [C1][C2][C9]; publish **Manifest v1**.
3. Start data plane [C3→C6]; run canary tenant with synthetic load.



4. Wire [C7] with dummy ACK/CT; run [C8] replay; enforce **digest equality = 100%** for canary.
5. Enable [CTX-2] Gate; tag rails; confirm SPV loop [C11]  $\rightleftharpoons$  [C7].
6. Roll pilots under governance canary; promote after N clean windows.

#### 4.7 SLOs & alerts (minimum)

- **Window close p95** (C4):  $\leq 15$  min; warn 20; page 30.
- **Transcript seal** (C6):  $\leq 10$  min after close.
- **Replay equality** (C8): 100% on canary; any mismatch pages.
- **Gate correctness** (CTX-2): no ALLOW unless (digest match  $\wedge$  acceptance met).
- **Proof freshness** (C7): alert T-15 min; quorum breach pages.
- **Recon completeness** (C11): provider receipt within freshness window.

#### 4.8 Security defaults

- mTLS everywhere; sign manifests, transcript roots, gate responses.
- Dual-control on governance; audit every signing event.
- No PII in transcripts; separate encrypted bucket for CT/KYC evidence.
- Strict RBAC + NetworkPolicies; Kafka ACLs + Postgres RLS.

---

## 5) Implementation plan & checklists

### 5.1 0–30 days (MVP)

- Stand up platform; implement [C3][C4][C5][C6]; basic [C7] with ACK/CT only; [CTX-2] gate; [C11] minimal SPV intake; [C10] metrics.
- Deliver **Interface #1** + **#2**; tag rails with ids.

### 5.2 30–60 days (Pilot)

- Add **Interface #6** SPV loop; harden [C7] reason codes; enable [C8] replay; Observability SLO dashboards; governance canary.

### 5.3 60–90 days (GA)

- Add streaming (#5) if needed; ERP adaptor (#3) for finance-led teams; SDKs (#4) for dev-led teams; on-chain SPV (#7) where applicable; ops copilot (#8) as a support layer.

### Readiness checklist (excerpt)

- Determinism tests (unordered/FP/wall-clock bans)
  - Replay = compute on canary windows
  - Transcript bucket: immutability + replication
  - Gate idempotency on {window\_id, output\_digest}
  - Proof freshness alerts; quorum dashboards
  - Recon round-trip verified against provider
- 

## 6) Failure modes & runbooks

- **Digest mismatch (C8)** → **HOLD** at gate; attach replay diff; likely causes: non-canonical input, shard drift, accidental FP; action: freeze promotion; re-ingest/repair; replay; resume.
  - **Stale proofs (C7)** → **HOLD**; ops copilot nudges owners; action: obtain fresh ACK/CT/SPV; gate re-authorizes deterministically.
  - **Watermark stall (C4)** → slow/late partition; action: inspect lag; apply grace policy; re-balance partitions.
  - **Recon variance (C11)** → mismatch totals; action: re-normalize provider file; open reason code; block delta cohort.
- 

## 7) Glossary (short)

- **ACK**: Finance approval (e.g., reserves OK).
- **CT**: Compliance/TAX/KYC/rights attestation.
- **SPV**: Simplified Proof of Validity (provider receipts or on-chain headers/inclusion proofs).
- **Transcript**: Content-addressed, tiered record set (inputs/checkpoints/outputs) sufficient to replay and verify equivalence.
- **Output digest**: Hash over canonical allocation records + trailer (watermark, fold order, manifest hash).
- **Watermark**: Monotone condition that closes a window across partitions.

---

# Appendices

## A) Reason-code catalog (starter)

- DIGEST\_MISMATCH, STALE\_PROOF, INSUFFICIENT\_QUORUM, MISSING\_ACK, MISSING\_CT, MISSING\_SPV, PROVIDER\_TOTALS\_MISMATCH, WATERMARK\_STALL, POLICY\_VERSION\_CONFLICT.

## B) Payload sketches (JSON)

### Authorize (CTX-2)

```
{
  "window_id": "W-2025-09-12-1",
  "output_digest": "hex...",
  "decisions": [
    {"principal_id": "P123", "decision": "ALLOW"},
    {"principal_id": "P124", "decision": "HOLD", "reason_code":
"STALE_PROOF"}
  ]
}
```

### Acceptance (C7)

```
{
  "window_id": "W-2025-09-12-1",
  "kind": "CT",
  "subject_id": "P123",
  "status": "VALID",
  "quorum": 2,
  "expires_at": "2025-09-13T00:00:00Z",
  "signature": "base64..."
}
```

## C) API surface (sketch)

- /ingest/events:POST (C3)
- /window/{id}/authorize:POST (CTX-2)
- /audit/transcript/{window}:GET • /audit/output-digest/{window}:GET (C6)
- /acceptance/{window}:GET, /acceptance/proof:POST (C7)
- /replay/{window}:POST (C8)
- /policy/propose|activate|rollback (C1)
- /governance/decision:POST (C9)